

✓ DATA LOADING

```
import numpy as np
import pandas as pd
import tensorflow as tf
import matplotlib.pyplot as plt
import nltk
import re
from sklearn.model_selection import train_test_split
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import (Dense, Conv1D, MaxPooling1D, LSTM, Bidirectional,
                                     Embedding, Dropout, SpatialDropout1D, GlobalMaxPooling1D)
from tensorflow.keras.regularizers import l2
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau

data = pd.read_csv('/content/customer_service_sentiment - meaningful_customer_service_sentiment.csv')
```

✓ DATA PROCESSING

```
nltk.download('punkt')
nltk.download('stopwords')
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize

def preprocess_text(text):
    text = text.lower()
    text = re.sub(r'\W', ' ', text)
    text = re.sub(r'\s+', ' ', text)
    tokens = word_tokenize(text)
    tokens = [word for word in tokens if word not in stopwords.words('english')]
    return ' '.join(tokens)

# preprocessing
data['processed_text'] = data['customer_message'].apply(preprocess_text)

#MAPPING sentiments
sentiment_mapping = {'Frustrated': 0, 'Neutral': 1, 'Satisfied': 2}
data['sentiment_label'] = data['sentiment_label'].map(sentiment_mapping)
```

✓ TEXT EMBEDDINGS

```

from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer

# Bag of Words (BoW)
bow_vectorizer = CountVectorizer(max_features=5000)
X_train_bow = bow_vectorizer.fit_transform(train_data['processed_text']).toarray()
X_val_bow = bow_vectorizer.transform(val_data['processed_text']).toarray()
X_test_bow = bow_vectorizer.transform(test_data['processed_text']).toarray()

# TF-IDF
tfidf_vectorizer = TfidfVectorizer(max_features=5000)
X_train_tfidf = tfidf_vectorizer.fit_transform(train_data['processed_text']).toarray()
X_val_tfidf = tfidf_vectorizer.transform(val_data['processed_text']).toarray()
X_test_tfidf = tfidf_vectorizer.transform(test_data['processed_text']).toarray()

print("Embedding completed for BoW and TF-IDF.")

```

MODEL TRAINING WITH DIFFERENT ARCHITECTURES

✓ COMPARATIVE ANALYSIS OF MODELS

```

cnn_model.save('best_cnn_model.h5')
print("Best CNN model saved to best_cnn_model.h5")

```

✓ MODEL TRAINING AND COMAPRATIVE ANALYSIS

```

dropout_rate = 0.65
learning_rate = 0.0002
l2_factor = 1e-4
batch_size = 32

# =====
# CNN Model
# =====
cnn_model = Sequential([
    Embedding(input_dim=max_words, output_dim=128),
    SpatialDropout1D(0.5), # Slightly reduced dropout on embeddings
    Conv1D(filters=32, kernel_size=5, activation='relu', kernel_regularizer=l2(l2_factor)),
    tf.keras.layers.BatchNormalization(),
    GlobalMaxPooling1D(),
    Dropout(dropout_rate),
    Dense(16, activation='relu', kernel_regularizer=l2(l2_factor)),
    tf.keras.layers.BatchNormalization(),
    Dropout(dropout_rate),
    Dense(3, activation='softmax')
])

```

[illegible]

verbose=1)

```
def evaluate_model(model, X, y, name):  
    loss, acc = model.evaluate(X, y, verbose=0)  
    print(f"{name} Model Accuracy on test data: {acc:.4f}")
```

```
print("\nEvaluating models on test data:")  
evaluate_model(cnn_model, X_test, y_test, "CNN")  
evaluate_model(lstm_model, X_test, y_test, "LSTM")  
evaluate_model(bilstm_model, X_test, y_test, "BiLSTM")
```



```
[nltk_data] Downloading package punkt to /root/nltk_data...  
[nltk_data] Package punkt is already up-to-date!  
[nltk_data] Downloading package stopwords to /root/nltk_data...  
[nltk_data] Package stopwords is already up-to-date!  
Training CNN Model...
```

Epoch 1/10

100/100 ————— 7s 13ms/step - accuracy: 0.3305 - loss: 1.7905 -

Epoch 2/10

100/100 ————— 1s 5ms/step - accuracy: 0.3878 - loss: 1.5052 -

Epoch 3/10

100/100 ————— 0s 5ms/step - accuracy: 0.4350 - loss: 1.3354 -

Epoch 4/10

100/100 ————— 0s 5ms/step - accuracy: 0.4882 - loss: 1.1513 -

Epoch 5/10

100/100 ————— 1s 6ms/step - accuracy: 0.5072 - loss: 1.0413 -

Epoch 6/10

100/100 ————— 0s 5ms/step - accuracy: 0.5554 - loss: 1.0045 -

Epoch 7/10

100/100 ————— 1s 5ms/step - accuracy: 0.5755 - loss: 0.8875 -

Epoch 8/10

100/100 ————— 1s 5ms/step - accuracy: 0.6047 - loss: 0.8518 -

Epoch 9/10

100/100 ————— 1s 5ms/step - accuracy: 0.6658 - loss: 0.7693 -

Epoch 10/10

100/100 ————— 1s 4ms/step - accuracy: 0.6685 - loss: 0.7221 -

Restoring model weights from the end of the best epoch: 10.

Training LSTM Model...

Epoch 1/12

100/100 ————— 31s 247ms/step - accuracy: 0.3481 - loss: 1.7180 -

Epoch 2/12

100/100 ————— 41s 249ms/step - accuracy: 0.3558 - loss: 1.6210 -

Epoch 3/12

100/100 ————— 39s 233ms/step - accuracy: 0.3630 - loss: 1.4880 -

Epoch 4/12

100/100 ————— 42s 240ms/step - accuracy: 0.3559 - loss: 1.4438 -

Epoch 5/12

100/100 ————— 42s 250ms/step - accuracy: 0.3593 - loss: 1.4314 -

Epoch 6/12

100/100 ————— 41s 247ms/step - accuracy: 0.3753 - loss: 1.3736 -

Epoch 7/12

100/100 ————— 40s 233ms/step - accuracy: 0.3962 - loss: 1.3077 -

Epoch 8/12

100/100 ————— 24s 236ms/step - accuracy: 0.3928 - loss: 1.2956 -

Epoch 9/12

```

100/100 ————— 25s 250ms/step - accuracy: 0.3990 - loss: 1.2820
Epoch 10/12
100/100 ————— 25s 248ms/step - accuracy: 0.4164 - loss: 1.2619
Epoch 11/12
100/100 ————— 41s 249ms/step - accuracy: 0.4261 - loss: 1.2356
Epoch 12/12
100/100 ————— 41s 252ms/step - accuracy: 0.4695 - loss: 1.2003
Restoring model weights from the end of the best epoch: 12.
Training BiLSTM Model...
Epoch 1/15
100/100 ————— 54s 433ms/step - accuracy: 0.3302 - loss: 2.0759
Epoch 2/15
100/100 ————— 42s 424ms/step - accuracy: 0.3318 - loss: 1.9273
Epoch 3/15

```

✓ SAVING THE BEST MODEL

```

cnn_model.save('best_cnn_model.h5')
print("Best CNN model saved to best_cnn_model.h5")

# predict sentiment for new text data
def predict_sentiment(text, model, tokenizer, max_length=100):
    # Preprocess the text using your existing function
    processed_text = preprocess_text(text)

    # Convert text to sequence and pad it
    sequence = tokenizer.texts_to_sequences([processed_text])
    padded_sequence = pad_sequences(sequence, maxlen=max_length)

    prediction = model.predict(padded_sequence)
    predicted_class = np.argmax(prediction, axis=1)[0]

    #mapping numeric prediction back to sentiment label
    sentiment_mapping = {0: 'Frustrated', 1: 'Neutral', 2: 'Satisfied'}
    predicted_sentiment = sentiment_mapping[predicted_class]

    return predicted_sentiment, prediction

:
new_text = "I am really happy with your service, everything was excellent."
sentiment, probabilities = predict_sentiment(new_text, cnn_model, tokenizer, max_length)
print("Predicted sentiment:", sentiment)
print("Prediction probabilities:", probabilities)

⇒ 1/1 ————— 1s 1s/step
Predicted sentiment: Satisfied
Prediction probabilities: [[0.24302487 0.2500682 0.50690687]]

```

